

16824 Homework 1

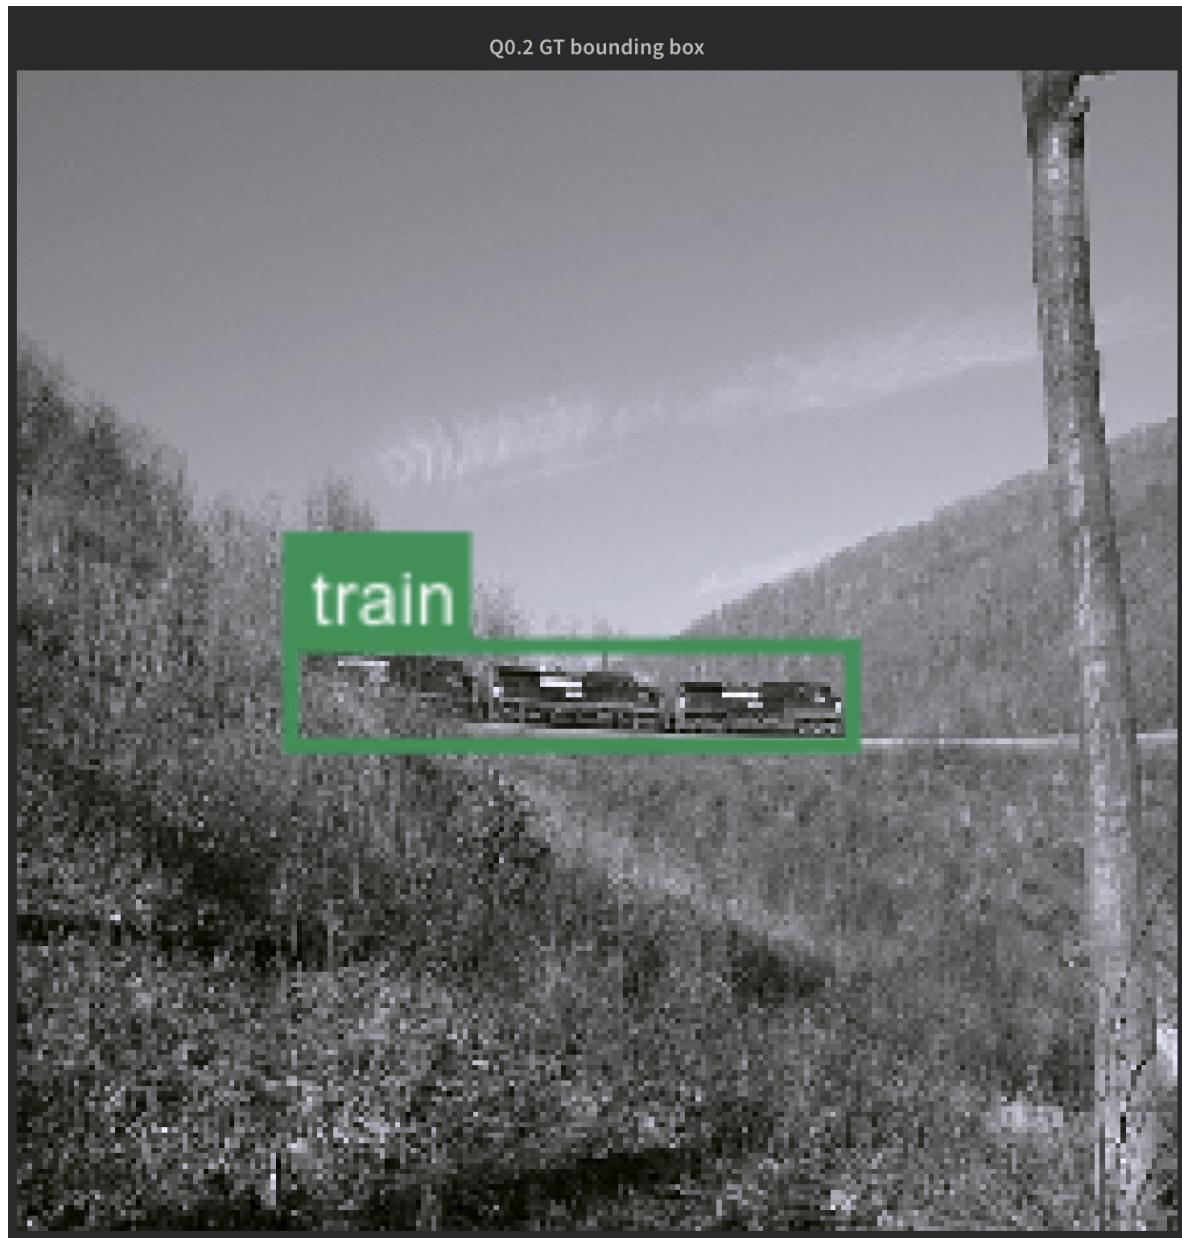
Jiajun Wan
Andrew ID: jiajunw2

Task 0: Visualization and Understanding the Data Structures

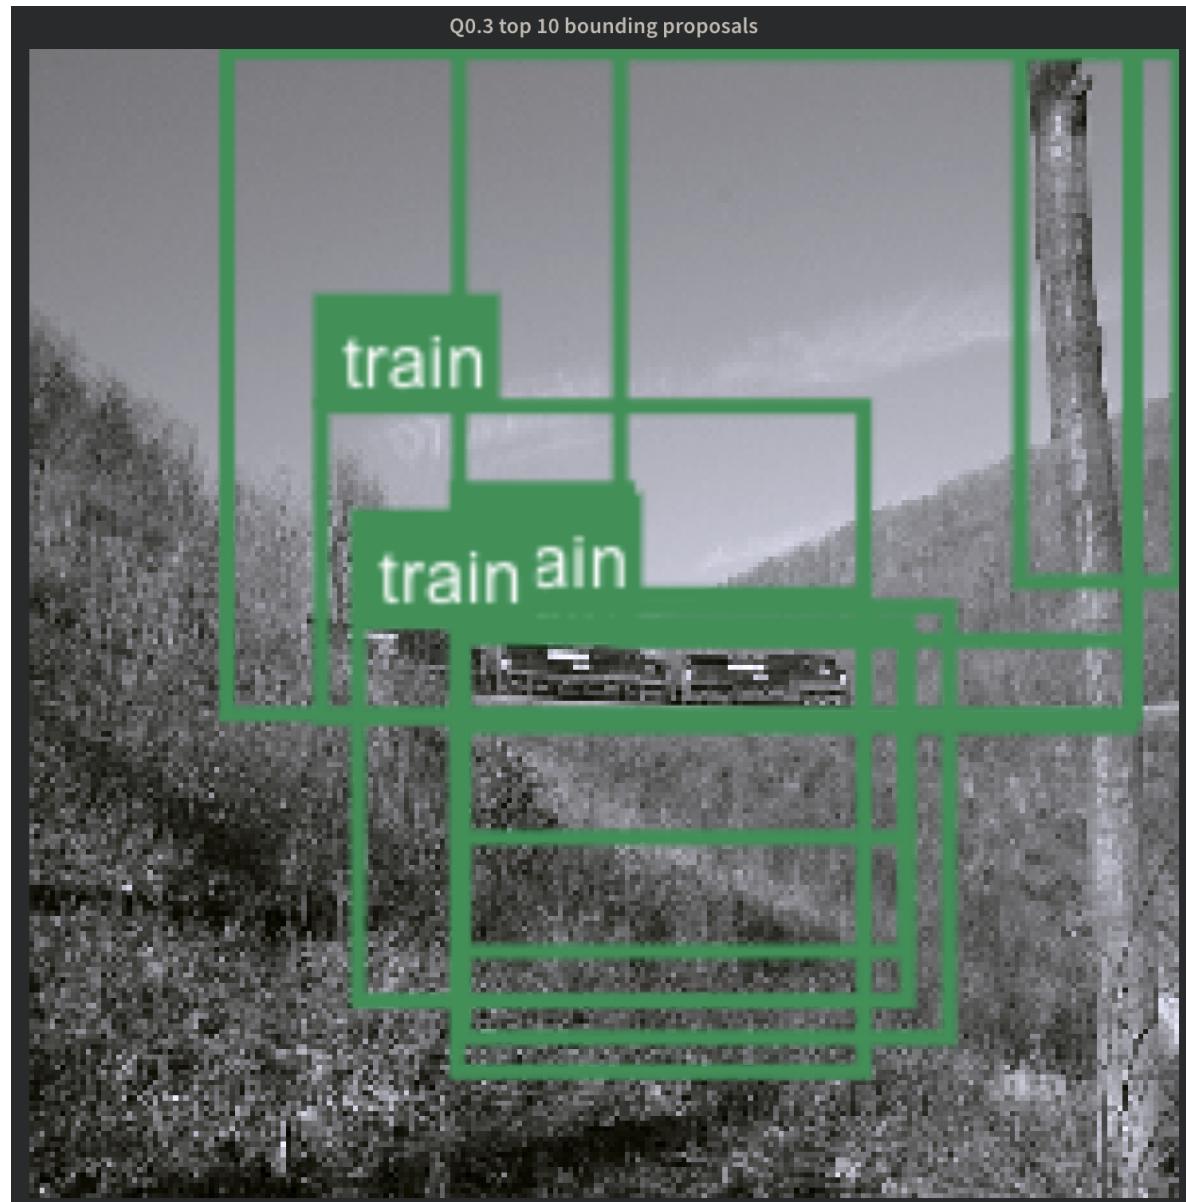
Q 0.1: What classes does the image at index 2020 contain (index 2020 is the 2021-th image due to 0-based numbering)?

It contains train.

Q 0.2 Use Wandb to visualize the ground-truth bounding box and the class for the image at index 2020.



Q 0.3 Use Wandb to visualize the top ten bounding box proposals for the image at index 2020.



Q1.1 describe functionality of the completed TODO blocks with comments

```
1 # TODO (Q1.1): define loss function (criterion) and optimizer from [1]
2 # also use an LR scheduler to decay LR by 10 every 30 epochs
3 # Binary Cross Entropy Loss
4 criterion = nn.BCEWithLogitsLoss()
5 # SGD
6 optimizer = torch.optim.SGD(model.parameters(),
7                             lr=args.lr,
8                             momentum=args.momentum,
9                             weight_decay=args.weight_decay)
10 # LR Step scheduler
11 scheduler = torch.optim.lr_scheduler.StepLR(optimizer,
12                                              step_size=30,
13                                              gamma=0.1)
14
15 # TODO (Q1.1): Create Datasets and Dataloaders using VOCdataset
16 # Ensure that the sizes are 512x512
17 # Also ensure that data directories are correct
18 # The ones use for testing by TAs might be different
19 # create train and validation dataset
20 train_dataset = VOCdataset(split='trainval',
21                            image_size=512)
22 val_dataset = VOCdataset(split='test',
23                         image_size=512)
24 train_sampler = None
25
26 # create train and validation dataloader with custom collate function
27 train_loader = torch.utils.data.DataLoader(
28     train_dataset,
29     batch_size=args.batch_size,
30     shuffle=False,
31     num_workers=args.workers,
32     collate_fn=train_dataset.collate_fn, # use custom collate function
33     pin_memory=True,
34     sampler=train_sampler,
35     drop_last=True)
36
37 val_loader = torch.utils.data.DataLoader(
38     val_dataset,
39     batch_size=args.batch_size,
40     shuffle=False,
41     num_workers=args.workers,
42     collate_fn=val_dataset.collate_fn, # use custom collate function
43     pin_memory=True,
44     drop_last=True)
45
46 # TODO (Q1.1): Get inputs from the data dict
47 # Convert inputs to cuda if training on GPU
48 # clear optimizer gradient
49 optimizer.zero_grad()
```

```

50 input = data[0].cuda()
51 target = data[1].cuda()
52
53 # TODO (Q1.1): Get output from model
54 output = model(input)
55
56 # TODO (Q1.1): Perform any necessary operations on the output
57
58 # TODO (Q1.1): Compute loss using ``criterion``
59 loss = criterion(output, target)
60 # sigmoid on each element to enforce probability from 0 to 1
61 output = torch.sigmoid(output)
62
63 # TODO (Q1.1): compute gradient and perform optimizer step
64 # backprop
65 loss.backward()
66 optimizer.step()
67
68 # TODO (Q1.1): Get inputs from the data dict
69 # Convert inputs to cuda if training on GPU
70 input = data[0].cuda()
71 target = data[1].cuda()
72
73 # TODO (Q1.1): Get output from model
74 output = model(input)
75
76 # TODO (Q1.1): Perform any necessary functions on the output
77
78 # TODO (Q1.1): Compute loss using ``criterion``
79 loss = criterion(output, target)
80 # sigmoid on each element to enforce probability from 0 to 1
81 output = torch.sigmoid(output)
82
83 # TODO (Q1.1): Define model
84 # AlexNet feature extractor layers
85 self.features = nn.Sequential(
86     nn.Conv2d(3, 64, kernel_size=(11, 11), stride=(4, 4), padding=(2, 2)),
87     nn.ReLU(),
88     nn.MaxPool2d(kernel_size=(3, 3), stride=(2, 2), dilation=(1, 1), ceil_mode=False
89     ),
89     nn.Conv2d(64, 192, kernel_size=(5, 5), stride=(1, 1), padding=(2, 2)),
90     nn.ReLU(),
91     nn.MaxPool2d(kernel_size=(3, 3), stride=(2, 2), dilation=(1, 1), ceil_mode=False
91     ),
92     nn.Conv2d(192, 384, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)),
93     nn.ReLU(),
94     nn.Conv2d(384, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)),
95     nn.ReLU(),
96     nn.Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1)),
97     nn.ReLU(),
98 )

```

```

99 # our own classifier layers
100 self.classifier = nn.Sequential(
101     nn.Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1)),
102     nn.ReLU(),
103     nn.Conv2d(256, 256, kernel_size=(1, 1), stride=(1, 1)),
104     nn.ReLU(),
105     nn.Conv2d(256, num_classes, kernel_size=(1, 1), stride=(1, 1)),
106 )
107 # final global max pooling layers
108 self.pool = nn.Sequential(
109     nn.AdaptiveMaxPool2d((1, 1)),
110     nn.Flatten(),
111 )
112
113
114 def forward(self, x):
115     # TODO (Q1.1): Define forward pass
116     # forward propagation
117     x = self.features(x)
118     x = self.classifier(x)
119     x = self.pool(x)
120     return x
121
122 """
123 TODO:
124     1. Load bounding box proposals for the index from self.roi_data. The proposals
125        ↪ are of the format:
126        [y_min, x_min, y_max, x_max] or [top left row, top left col, bottom right row,
127        ↪ bottom right col]
128     2. Normalize in the range (0, 1) according to image size (be careful of width/
129        ↪ height and x/y correspondences)
130     3. Make sure to return only the top_n proposals based on proposal confidence ("'
131        ↪ boxScores")!
132     4. You may have to write a custom collate_fn since some of the attributes below
133        ↪ may be variable in number for each data point
134 """
135 # get array of (boxScore, boxes) tuples and sort by boxScore
136 score_bbox_pairs = sorted(list(zip(self.roi_data['boxScores'][0][index][:, 0], self.
137     ↪ roi_data['boxes'][0][index])), key=lambda x: x[0], reverse=True)[:self.top_n]
138 # convert proposal to [xmin, ymin, xmax, ymax], same as gt_boxes
139 proposals = np.asarray(list(map(lambda x: np.array([x[1][1] / width, x[1][0] /
140     ↪ height, x[1][3] / width, x[1][2] / height]), score_bbox_pairs))).astype(np.
141     ↪ float32)
142
143 def collate_fn(self, dict):
144     """
145         :param dict: a batch list of ret dictionary from __getitem__
146         :return: pytorch tensor of x and y
147     """
148
149     # extract input image and target label from dictionary and convert to pytorch
150     ↪ single precision float tensor

```

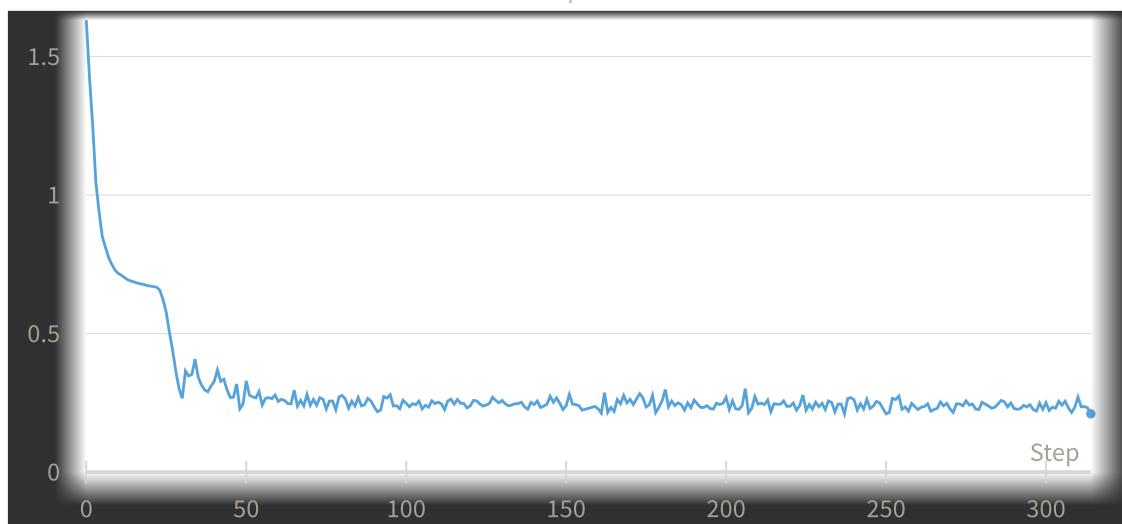
```
141 X = [x['image'] for x in dict]
142 Y = [y['label'] for y in dict]
143 X = torch.stack(X).to(torch.float32)
144 Y = torch.stack(Y).to(torch.float32)
145 return X, Y
```

Q 1.2 What is the shape ($N \times C \times H \times W$) of the output of the model?

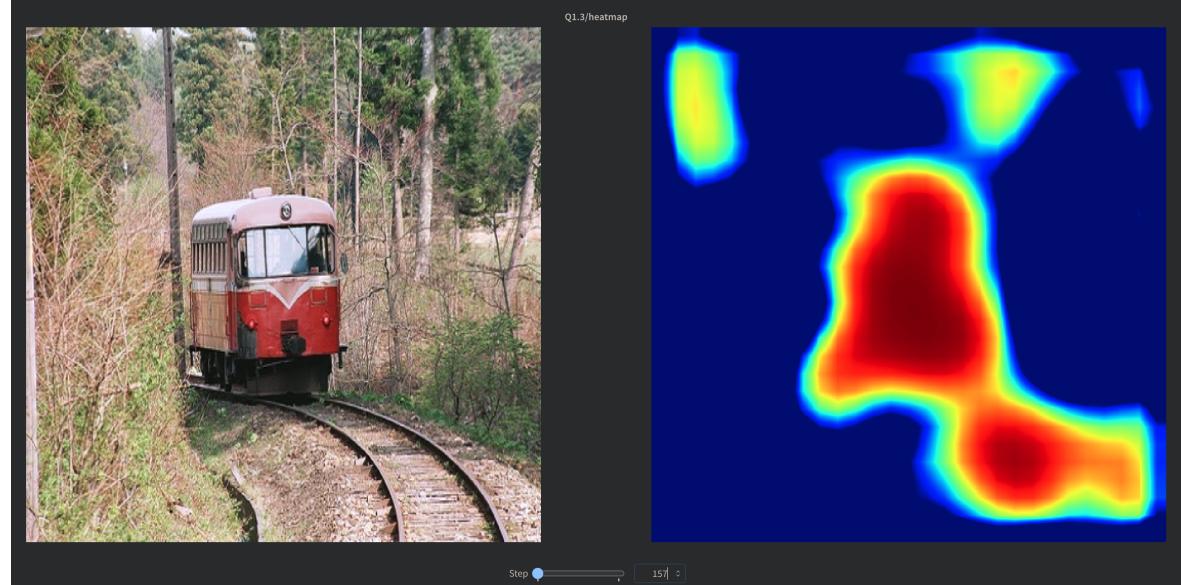
[Batch Size, 20, 1, 1]. After flatten: [Batch Size, 20]

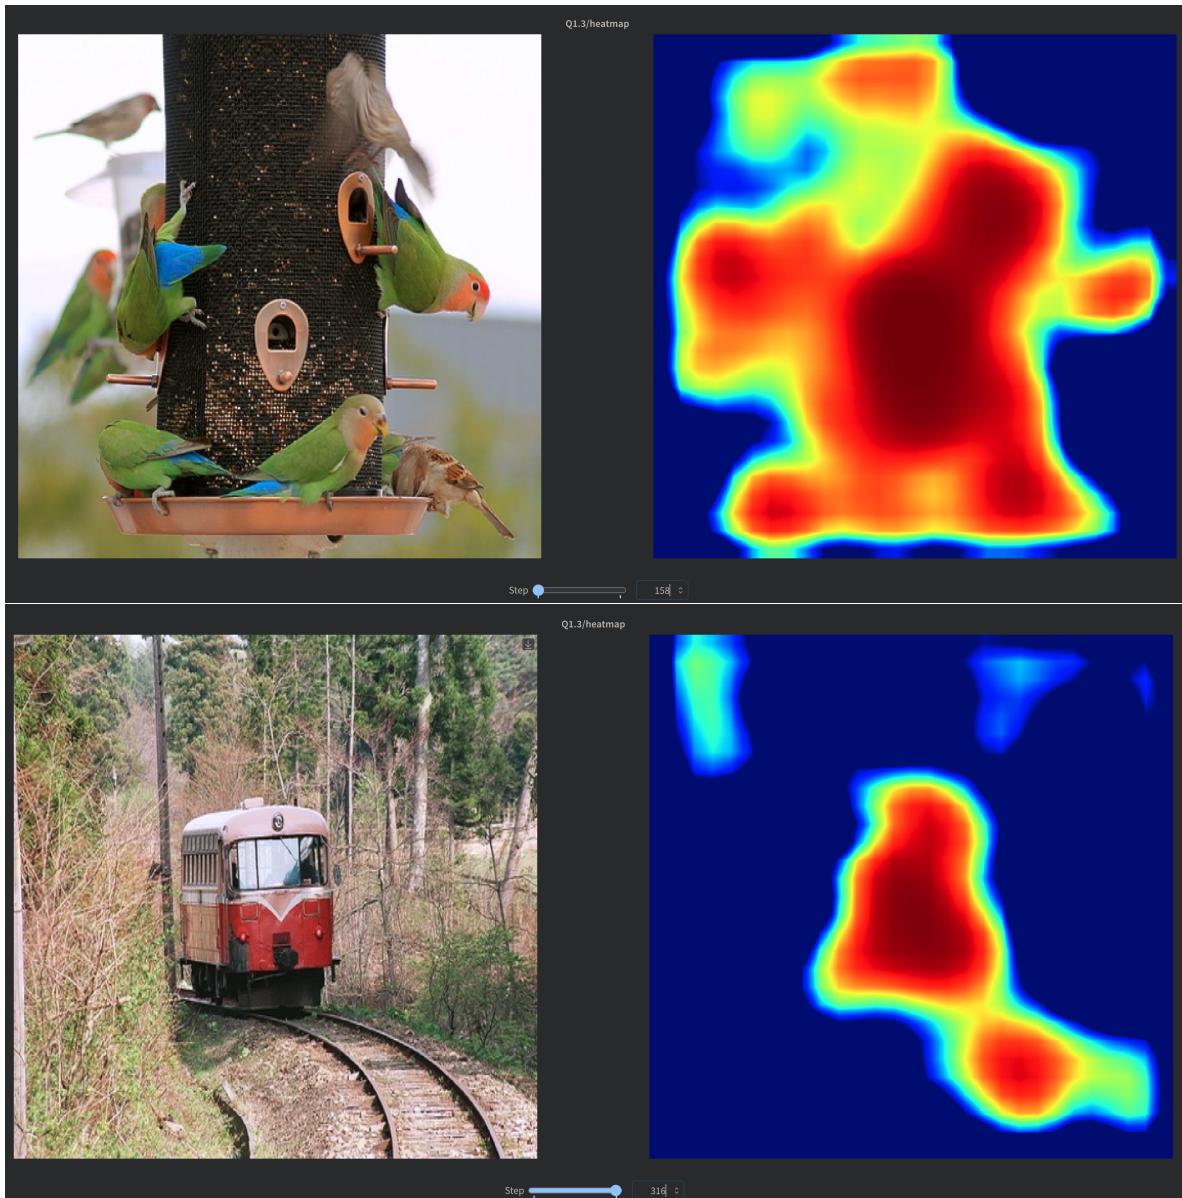
Q 1.3

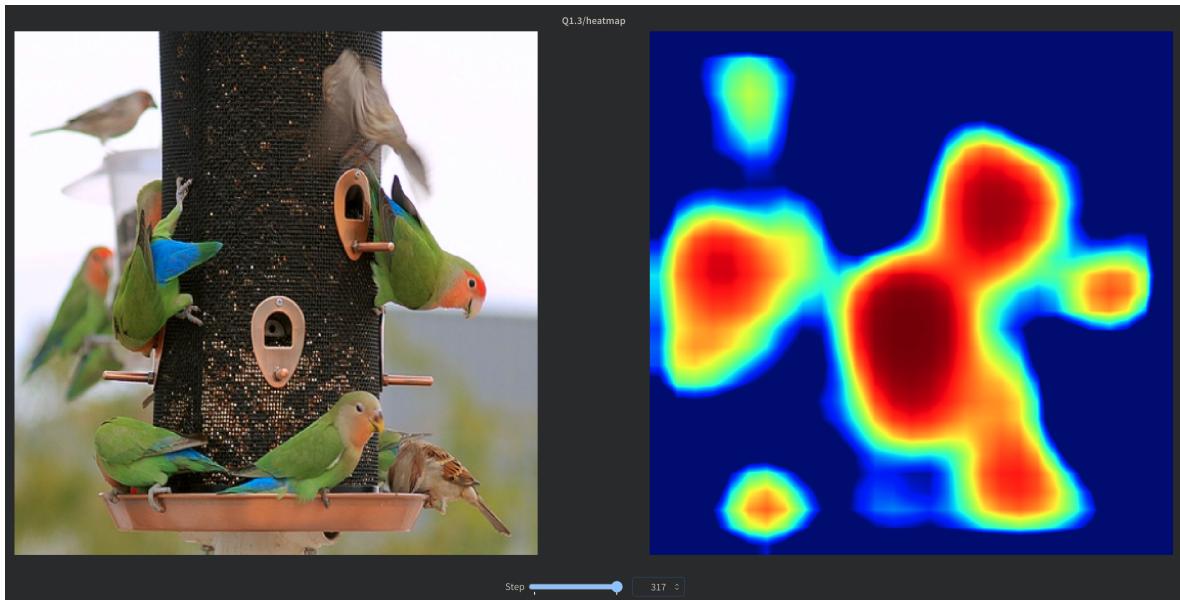
train/loss



Q1.3/heatmap







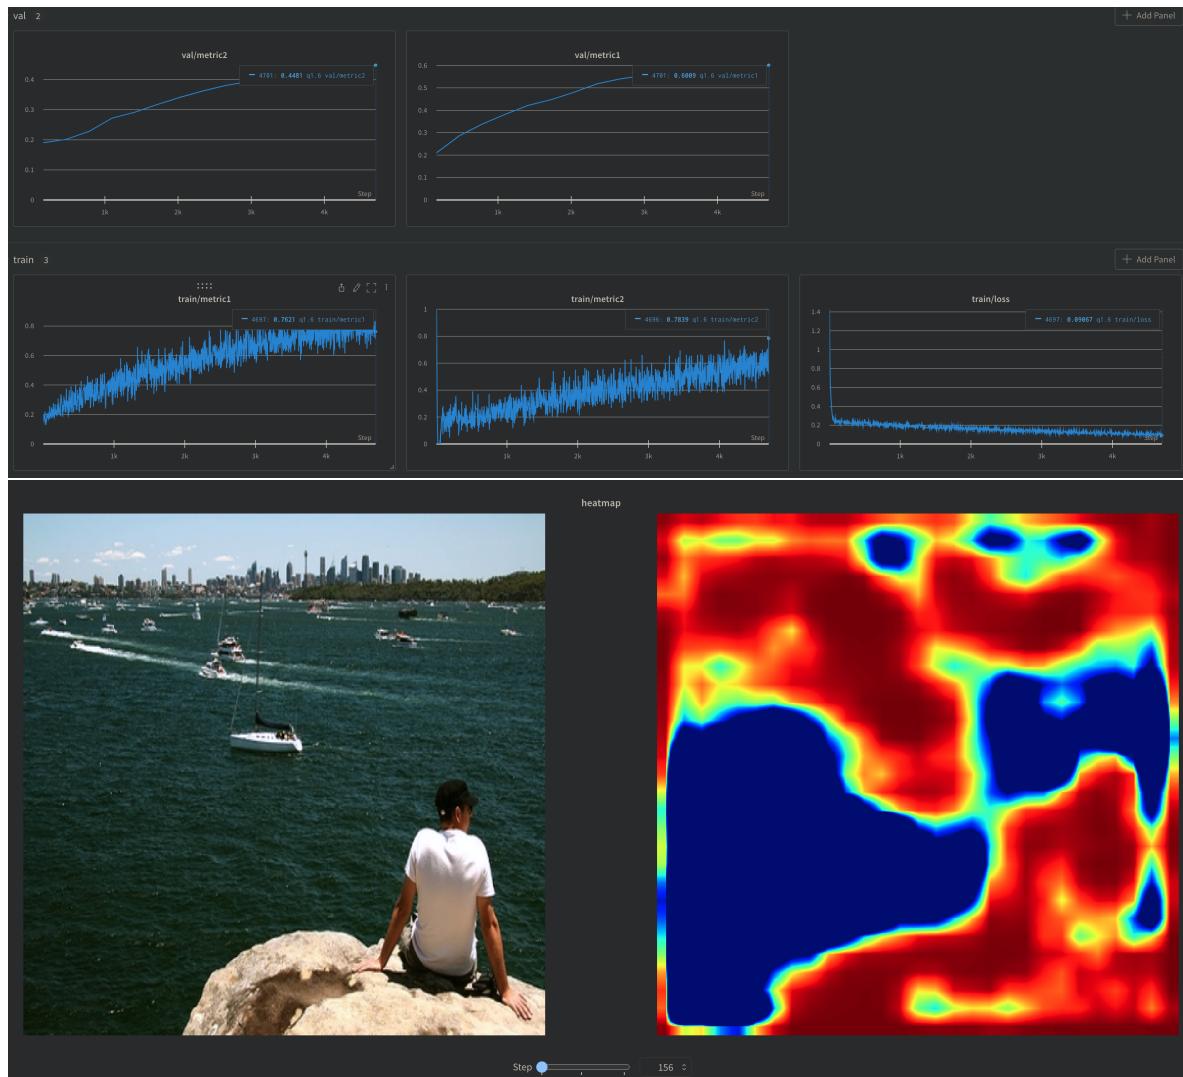
Q 1.4 In the first few iterations, you should observe a steep drop in the loss value. Why does this happen?

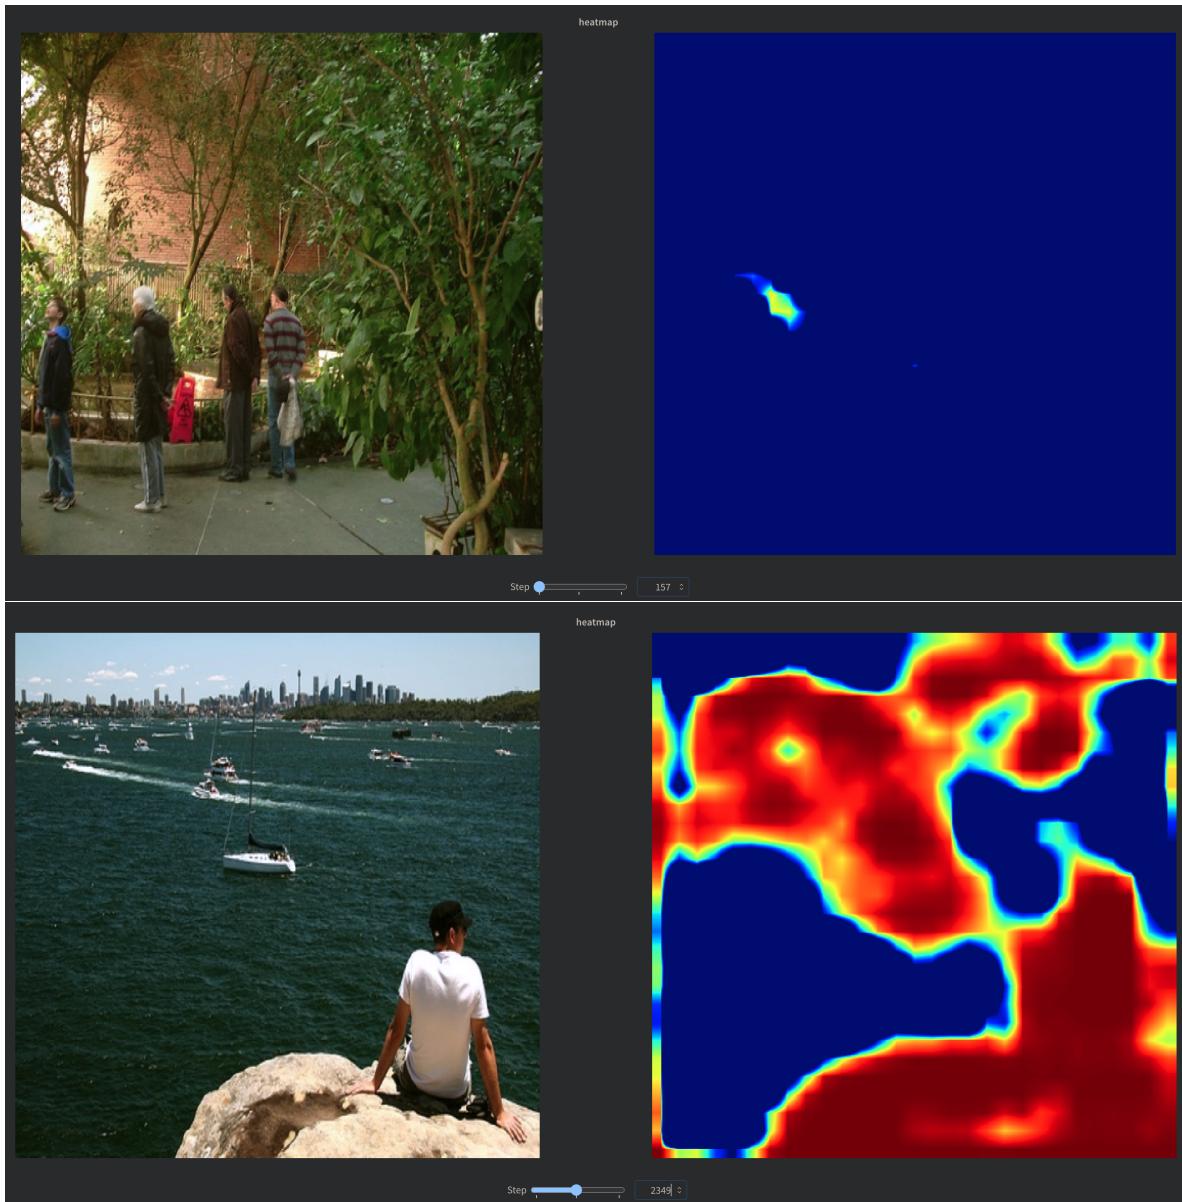
Because at first the network is initialized randomly, and there are multiple labels associated with an image, the loss is going to be large. Also the labels across a batch are not balanced with some labels in most images and some labels not appearing in any image. After the network is learning how to classify based on the activation heat maps, the loss decreases steeply. The loss gradually saturates after the network learns the most in the first few iterations.

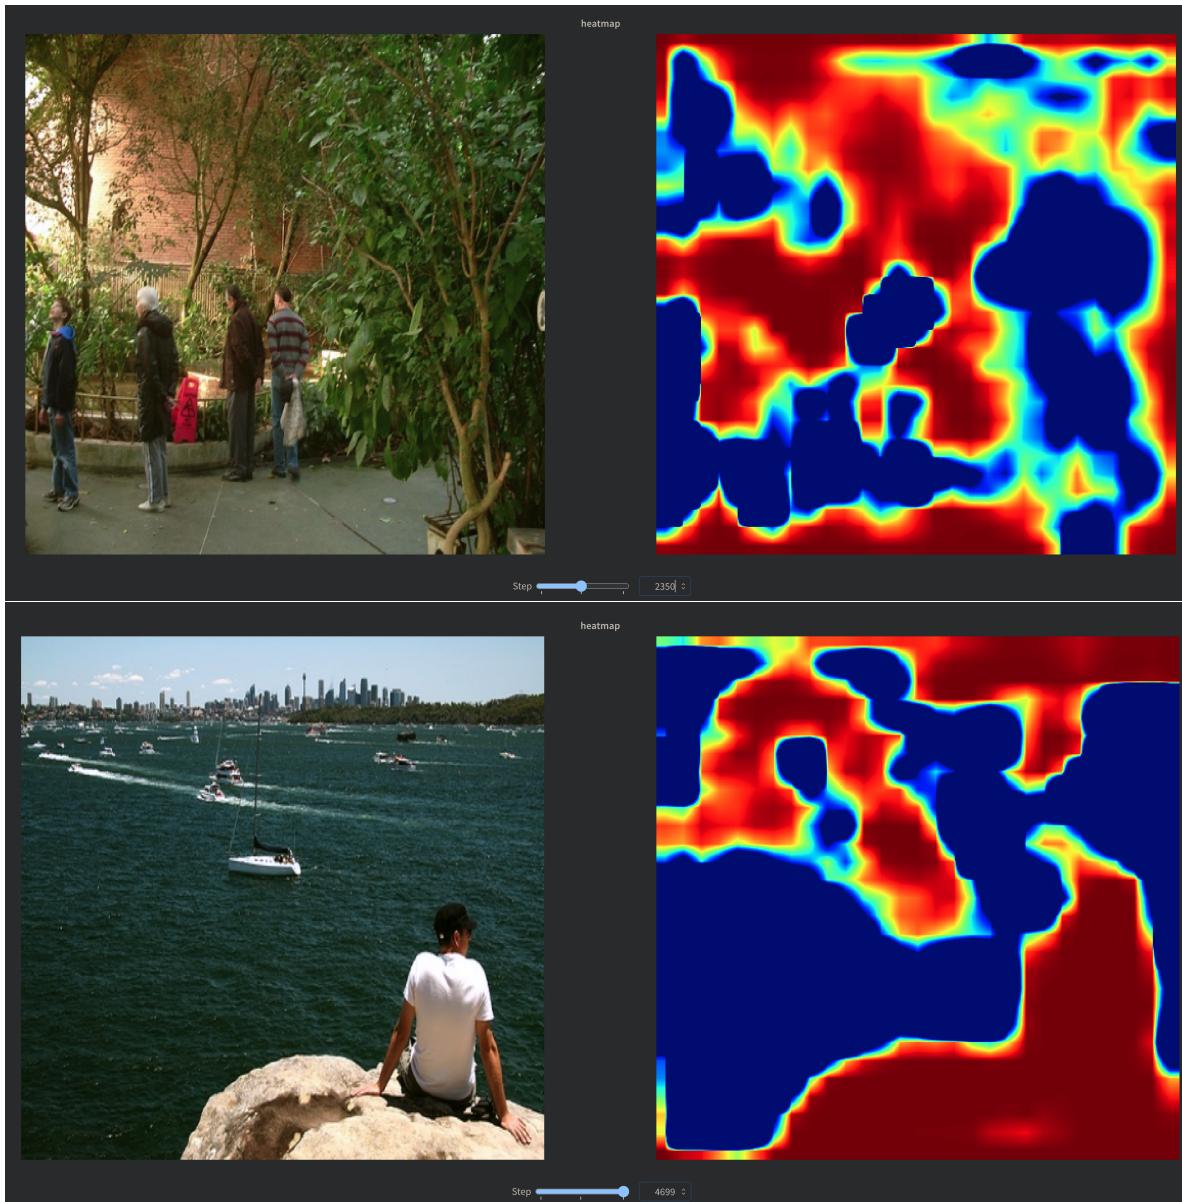
Q 1.5

Filled out TODOs in the code files. The assumption is threshold for deciding positive prediction is 0.5.

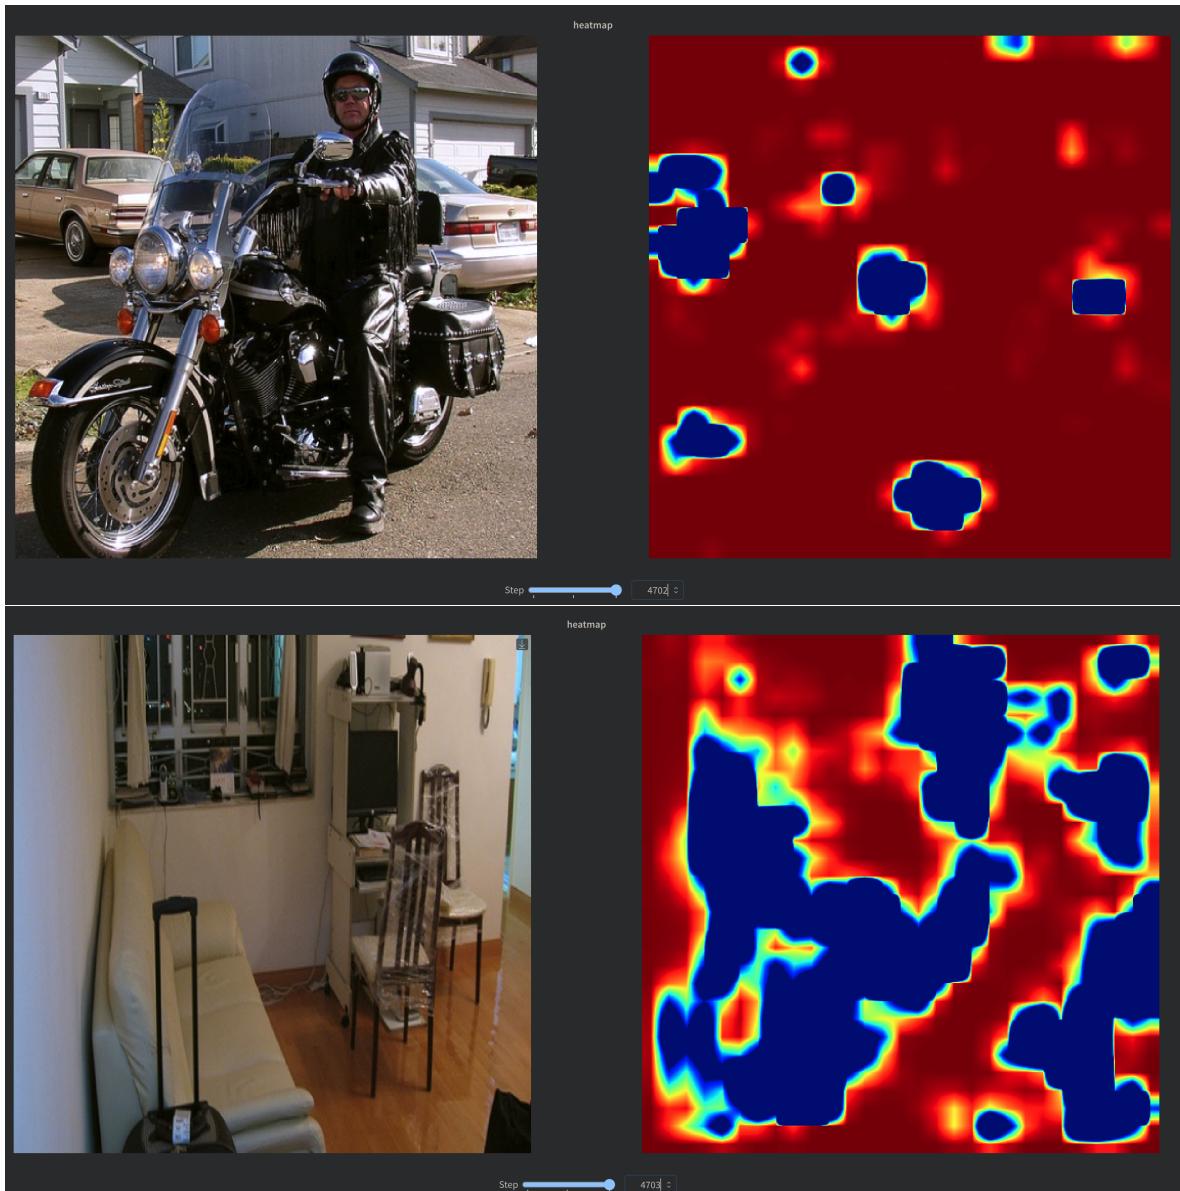
Q 1.6







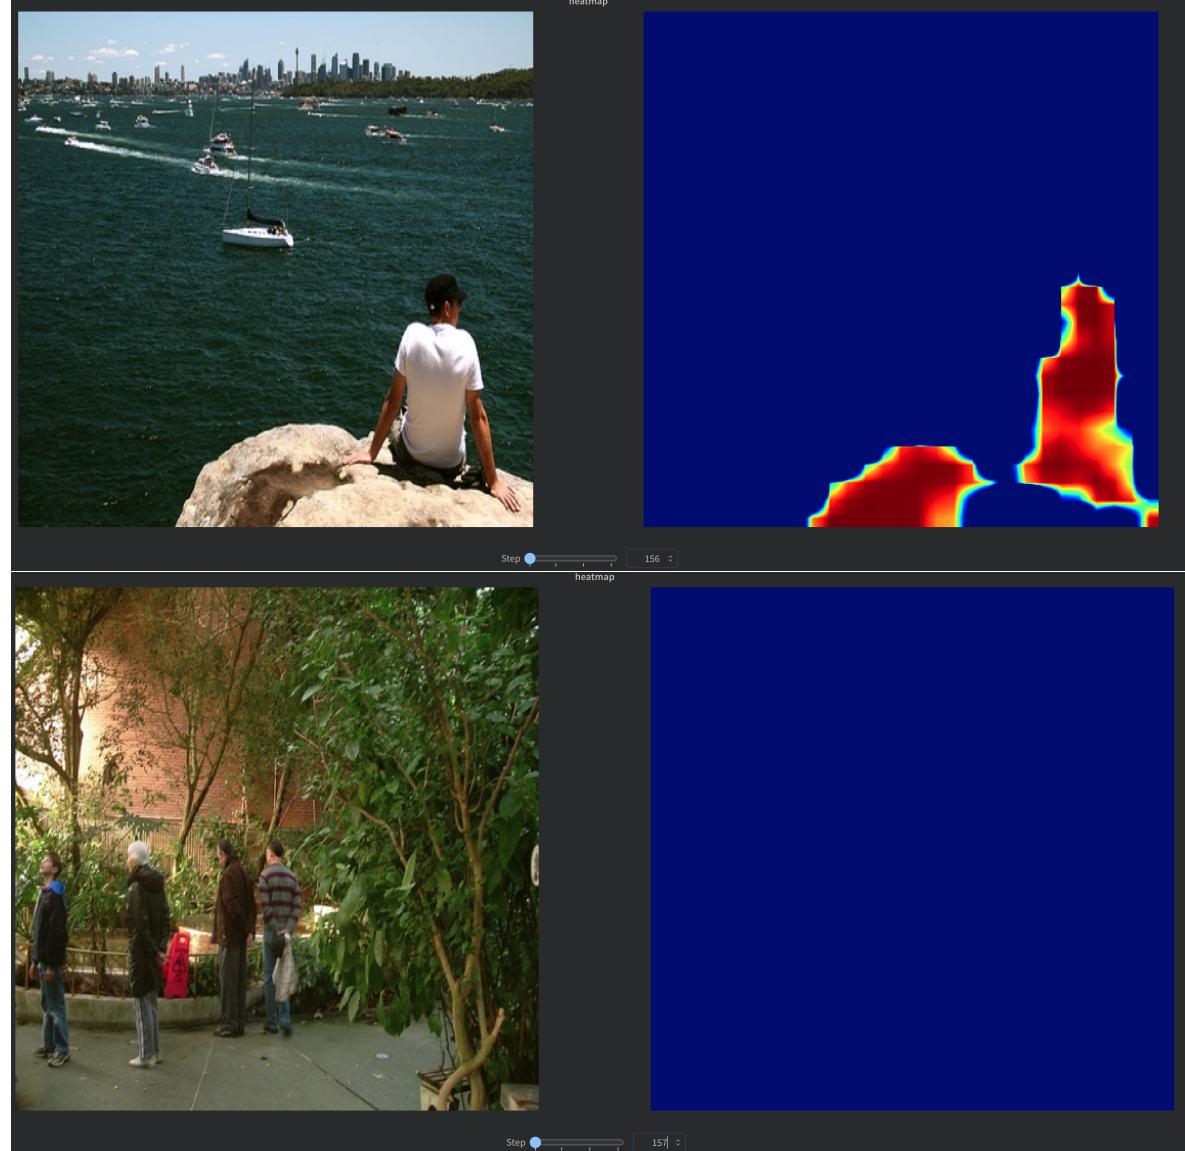


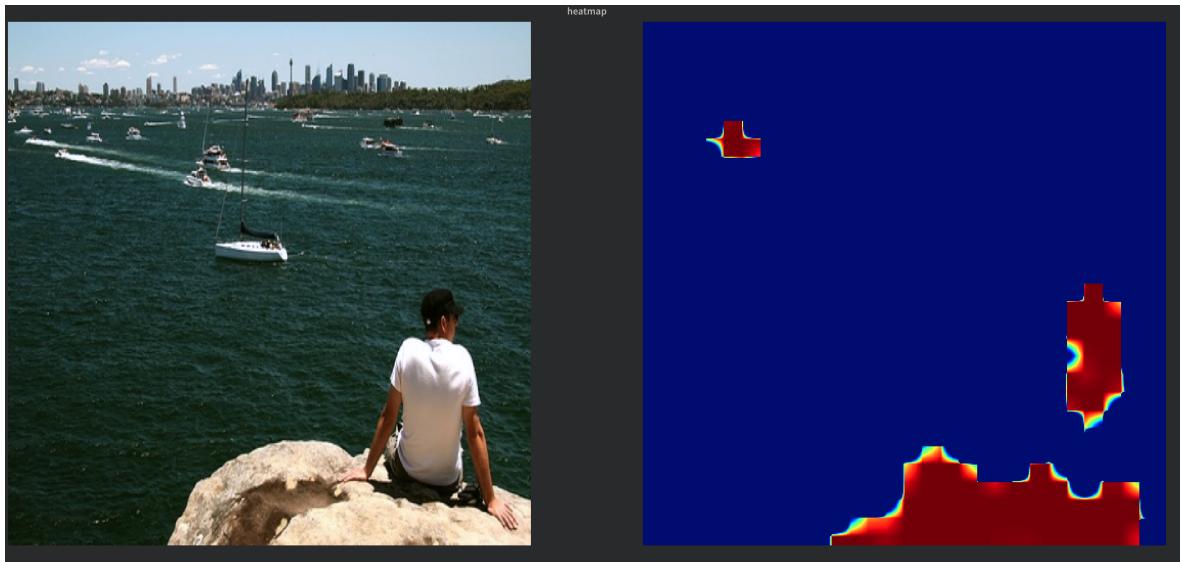


Training loss: 0.09067, metric1: 0.7621, metric2: 0.7839
Validation metric1: 0.4481, metric2: 0.6009

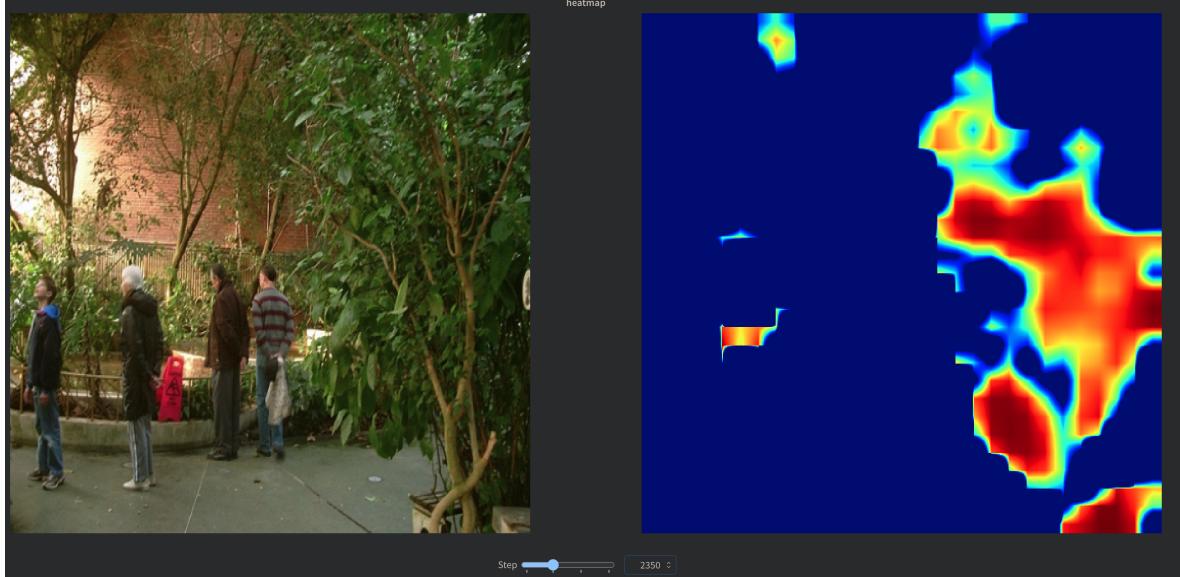
Q 1.7

I use global average pooling at the end to incorporate more salient features.

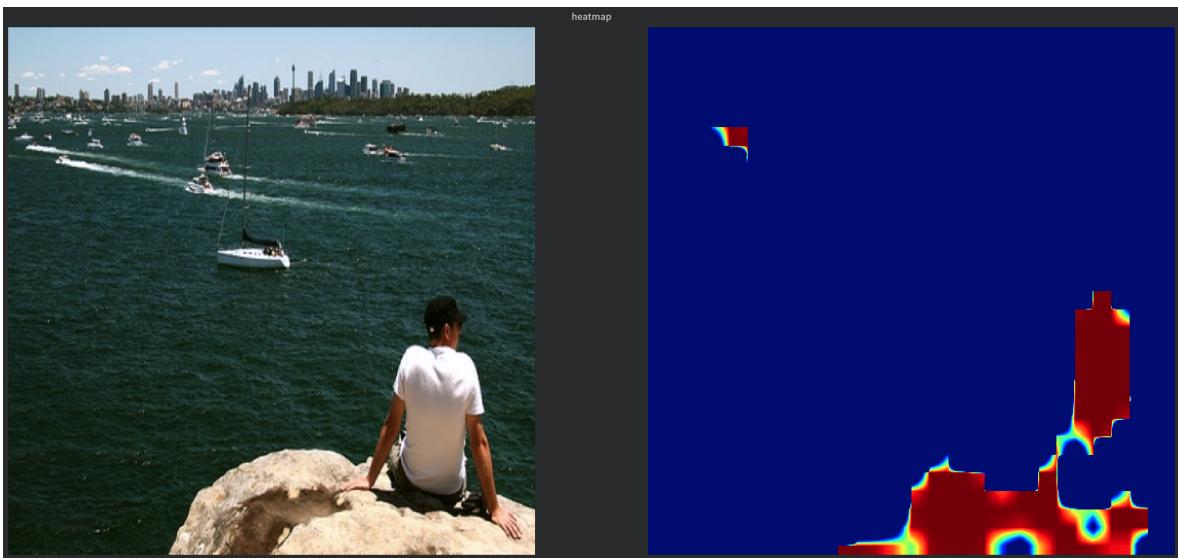




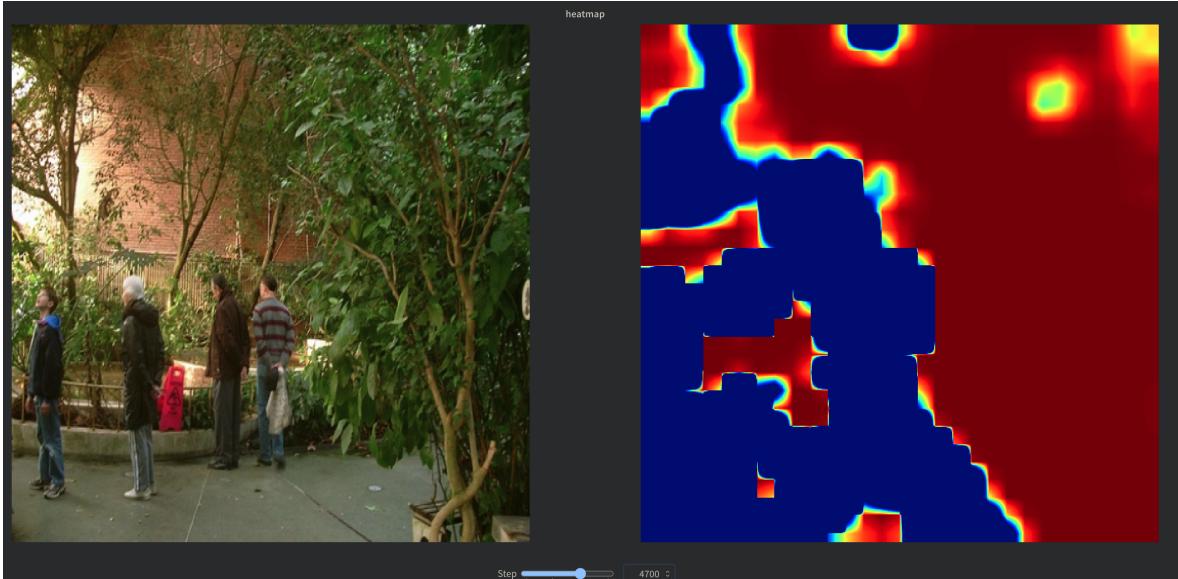
Step 2349



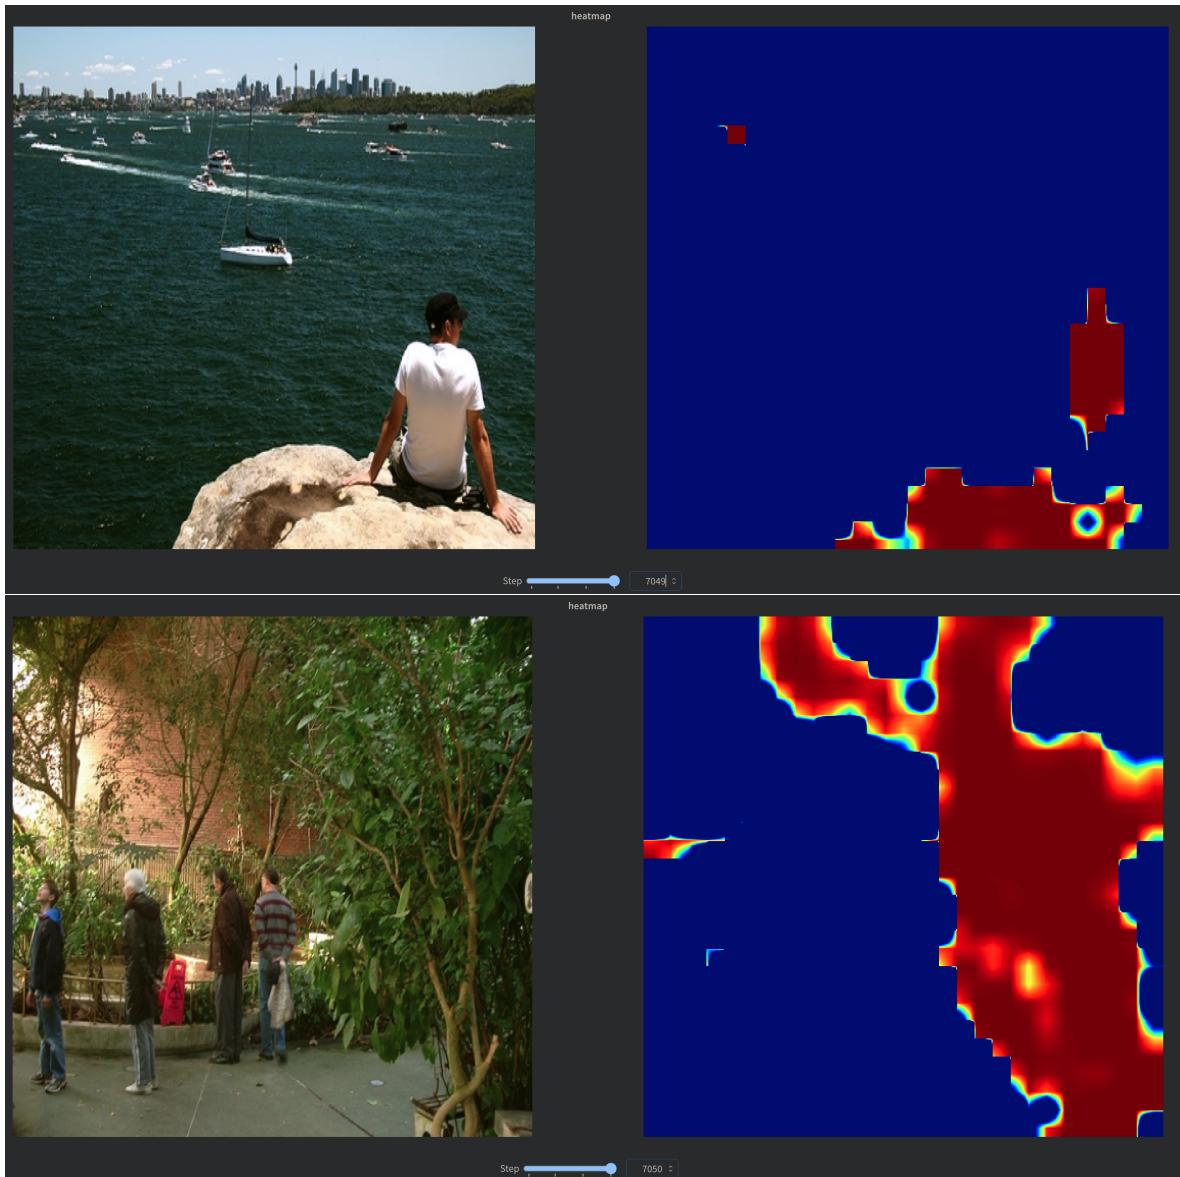
Step 2350



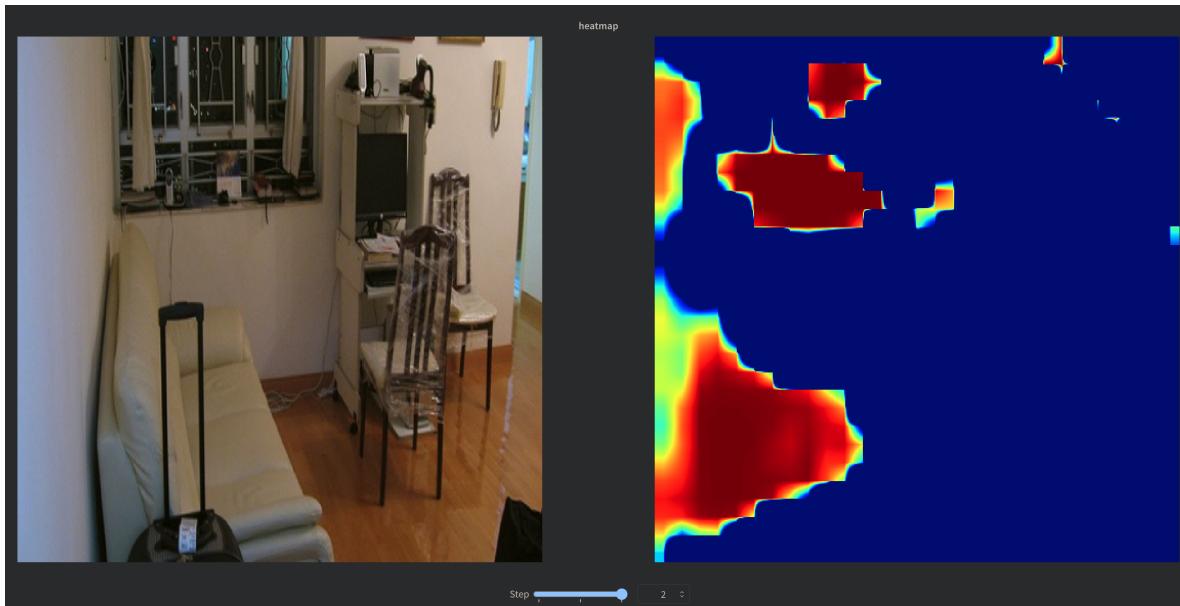
Step 4699



Step 4700







Training loss: 0.02455, metric1: 0.9821, metric2: 0.9453

Validation metric1: 0.7845, metric2: 0.7181

Q2.3

```
1 # TODO: given bounding boxes and corresponding scores, perform non max suppression
2 def nms(bounding_boxes, confidence_score, threshold=0.05):
3     """
4     bounding boxes of shape      Nx4
5     confidence scores of shape  N
6     threshold: confidence threshold for boxes to be considered
7
8     return: list of bounding boxes and scores
9     """
10    # sort bounding boxes based on confidence score
11    pairs = [x for x in list(zip(confidence_score, bounding_boxes)) if x[0] >
12              threshold]
13    pairs = sorted(pairs, key=lambda x: x[0], reverse=True)
14    boxes, scores = [], []
15    visited = set()
16    # loop to visit all bounding boxes
17    while len(visited) < len(pairs):
18        idx = 0
19        # get the first index that is not visited before
20        for i in range(len(pairs)):
21            if i in visited:
22                continue
23            idx = i
24            break
25        # add this bounding box to output list
26        boxes.append(pairs[idx][1])
27        scores.append(pairs[idx][0])
28        visited.add(idx)
29        # loop over other bounding boxes to remove overlapping bboxes
30        for i in range(len(pairs)):
31            if i in visited:
32                continue
33            # ignore overlapping bboxes with iou > 0.3
34            if iou(pairs[idx][1], pairs[i][1]) > 0.3:
35                visited.add(i)
36
37    return boxes, scores
38
39 # TODO: calculate the intersection over union of two boxes
40 def iou(box1, box2):
41     """
42     Calculates Intersection over Union for two bounding boxes (xmin, ymin, xmax,
43     ymax)
44     returns IoU value
45     """
46
47     # get the boarder coordiantes of union of two bboxes
48     x_left = max(box1[0], box2[0])
49     y_top = max(box1[1], box2[1])
50     x_right = min(box1[2], box2[2])
```

```

48     y_bottom = min(box1[3], box2[3])
49
50     # area is zero if two bboxes cannot be unioned
51     if x_right < x_left or y_bottom < y_top:
52         return 0.0
53
54     # calculate intersection area
55     intersection = (x_right - x_left) * (y_bottom - y_top)
56     area1 = (box1[2] - box1[0]) * (box1[3] - box1[1])
57     area2 = (box2[2] - box2[0]) * (box2[3] - box2[1])
58     # iou is intersection divided by union (addition of areas of two bboxes minus
59     # intersection)
60     iou = intersection / (area1 + area2 - intersection)
61     return iou
62
63 res = [[], []]
64 # TODO (Q2.3): Iterate over each class (follow comments)
65 labels = []
66 bboxes = []
67 for class_num in range(20):
68     # get valid rois and cls_scores based on thresh
69     # use NMS to get boxes and scores
70     boxes, scores = nms(rois.cpu(), cls_prob[:, class_num].cpu(), thresh)
71     if boxes and scores:
72         for i in range(len(boxes)):
73             res[0].append((scores[i], boxes[i], class_num, iter))
74             if args.use_wandb and (epoch == 0 or epoch == 4) and (cnt < 10 or iter
75             in logged):
76                 labels.append(class_num)
77                 bboxes.append(boxes[i].detach().numpy().astype(float))
78     res[1].append([iter, gt_boxes, gt_class_list])
79
80 def calculate_map(res):
81     """
82     Calculate the mAP for classification.
83     """
84
85     # TODO (Q2.3): Calculate mAP on test set.
86     # Feel free to write necessary function parameters.
87     # get predictions and ground truths
88     pred = res[0]
89     pred = sorted(pred, key=lambda x: x[0], reverse=True)
90     gt = res[1]
91     # iterate all ground truths to get number of gt_bboxes for each class
92     allgtboxes = [0] * 20
93     for i in range(len(gt)):
94         for j in range(len(gt[i][2])):
95             allgtboxes[gt[i][2][j]] += 1
96     # total for each class
97     AP = [0] * 20
98     TP = [0] * 20
99     FP = [0] * 20

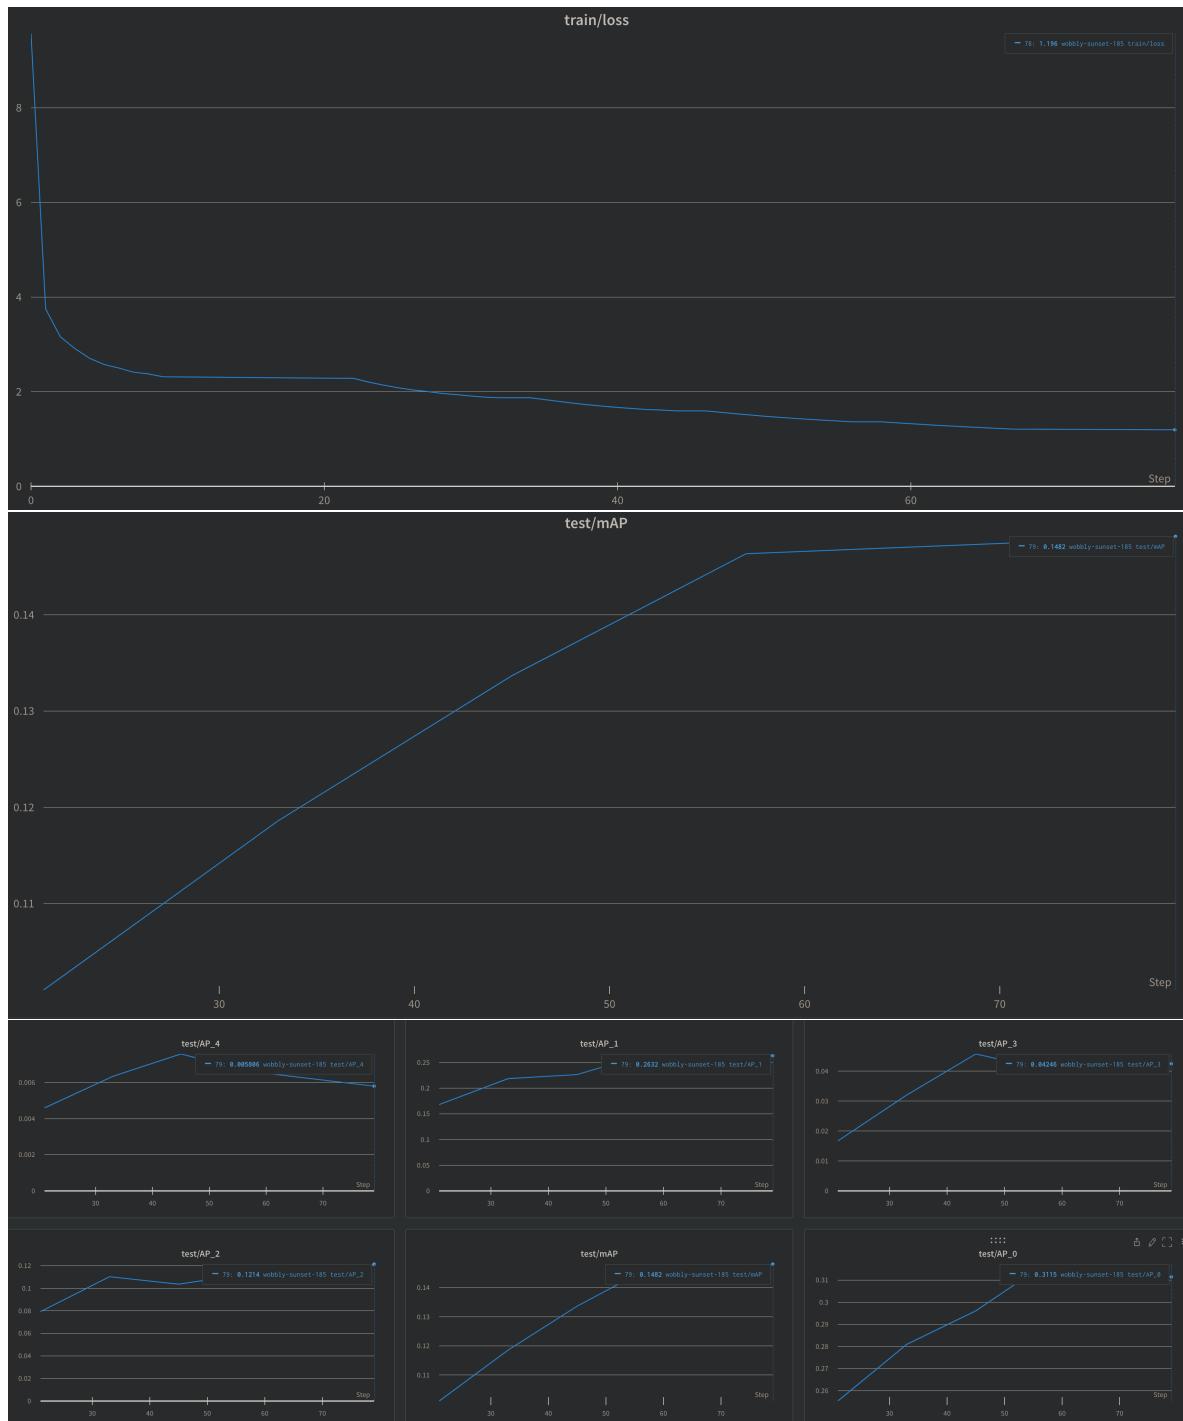
```

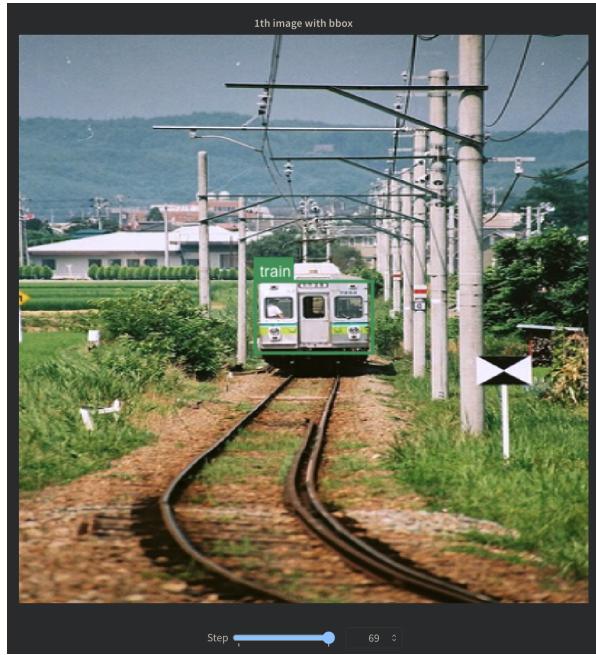
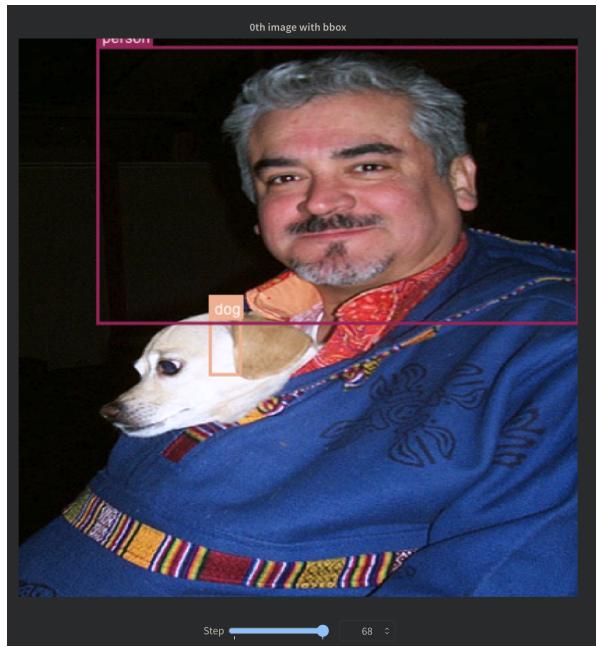
```

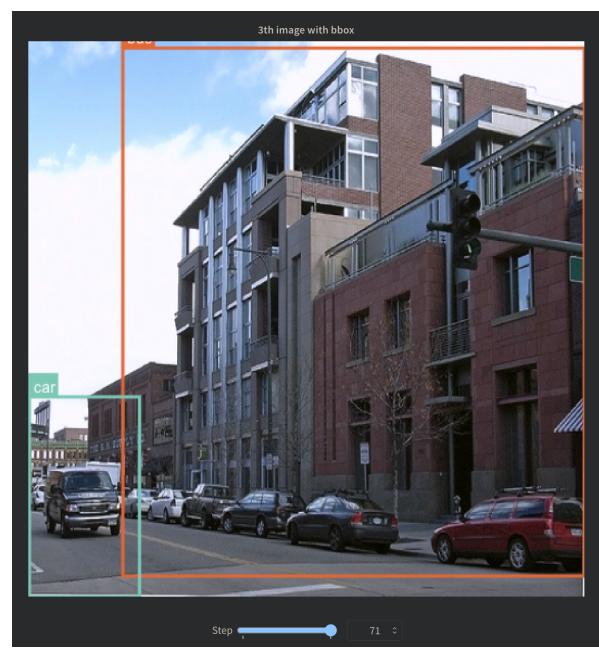
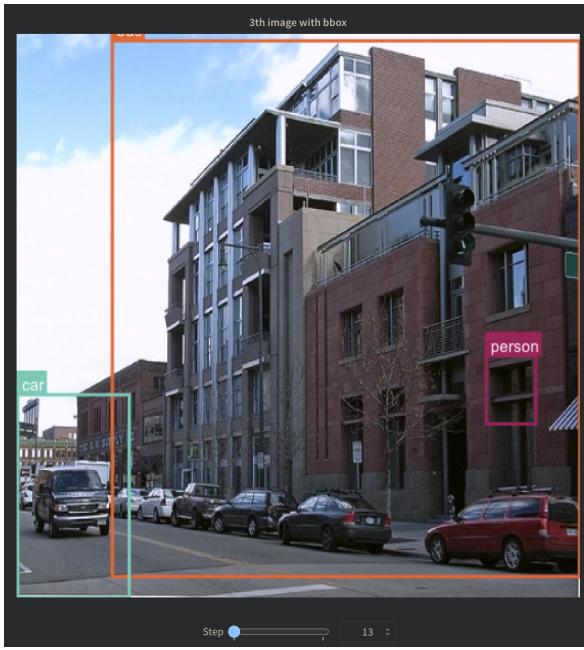
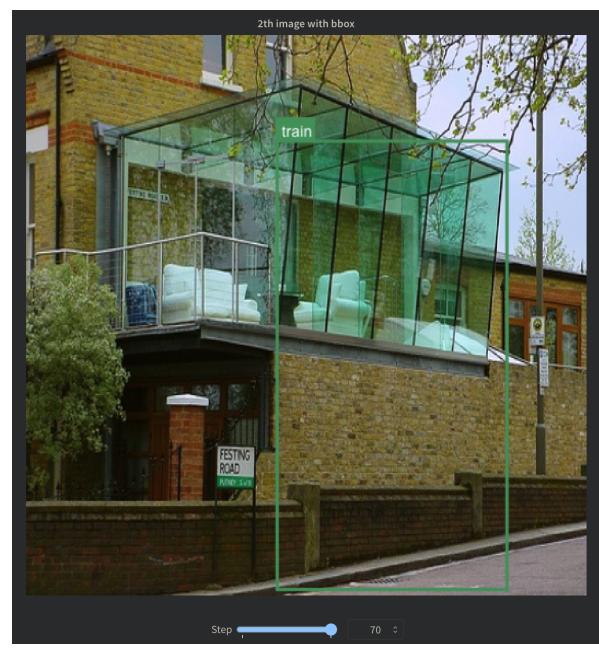
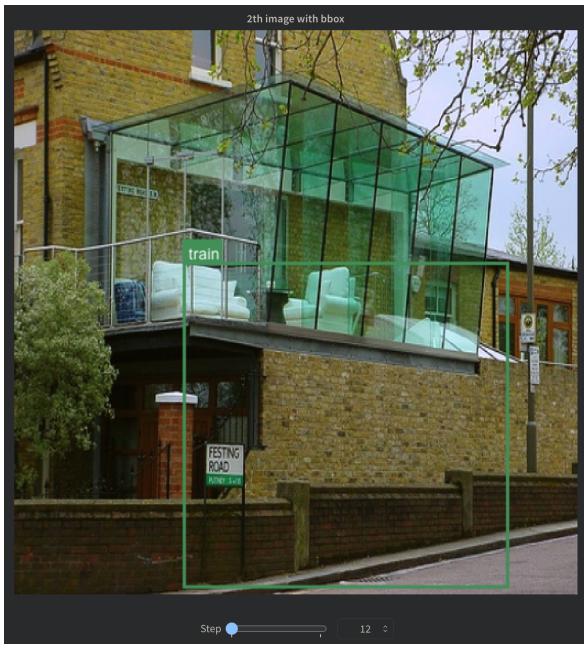
97     # visited set for each ground truth bboxes list, used for if one gt_bbox is
98     # it cannot be matched again
99     visited = [set() for _ in range(len(gt))]
100    # precision and recall list for AP (AUC) calculation for each class
101    precision = [[] for _ in range(20)]
102    recall = [[] for _ in range(20)]
103    # iterate prediction bounding boxes to calculate precision and recall
104    for _, box, class_num, iter in pred:
105        gt_boxes = gt[iter][1]
106        gt_class_list = gt[iter][2]
107        assert gt[iter][0] == iter
108        assert len(gt_boxes) == len(gt_class_list)
109        # if not more gt bboxes to match a prediction, this prediction is a false
110        # positive
111        if len(gt_boxes) == 0 or len(gt_boxes) == len(visited[iter]):
112            FP[class_num] += 1
113        else:
114            match = False
115            # iterate gt bboxes of the image
116            for i in range(len(gt_boxes)):
117                if i in visited[iter]:
118                    continue
119                # prediction is true positive if label is matched and iou exceeds
120                # threshold
121                if class_num == gt_class_list[i] and iou(box, gt_boxes[i]) > 0.45:
122                    TP[class_num] += 1
123                    visited[iter].add(i)
124                    match = True
125                    break
126                # prediction is not matched to any gt bboxes and is regarded a a false
127                # positive
128                if not match:
129                    FP[class_num] += 1
130
131            # add precision and recall to list
132            precision[class_num].append(TP[class_num] / (TP[class_num] + FP[class_num]))
133            recall[class_num].append(TP[class_num] / allgtbboxes[class_num])
134    # using sklearn auc function to calculate AP for each class
135    for i in range(20):
136        AP[i] = sklearn.metrics.auc(recall[i], precision[i])
137    return AP

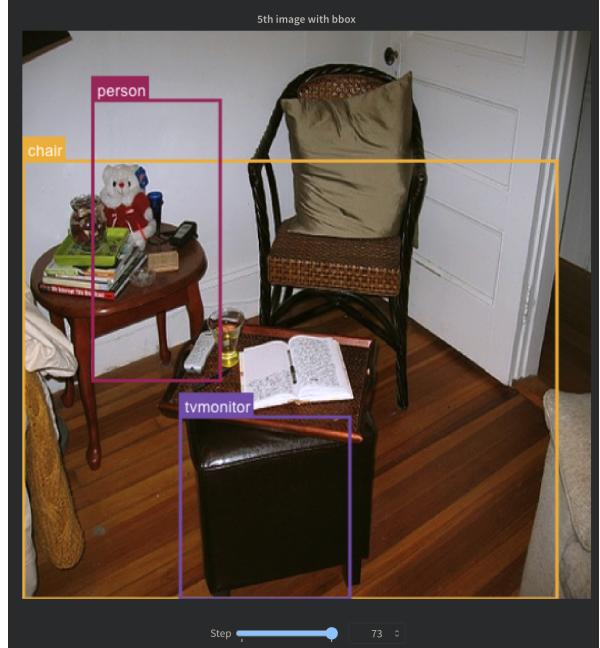
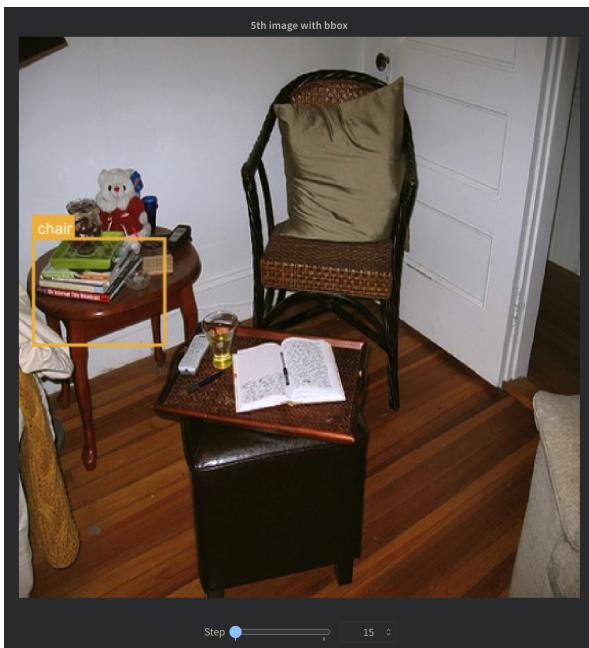
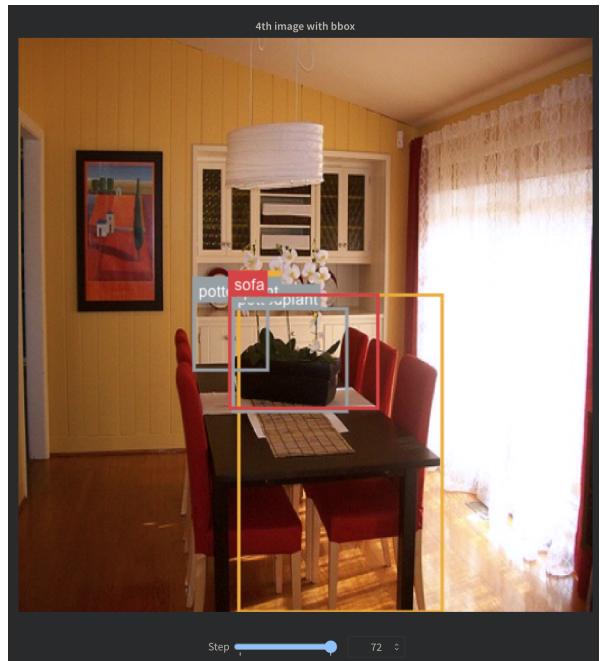
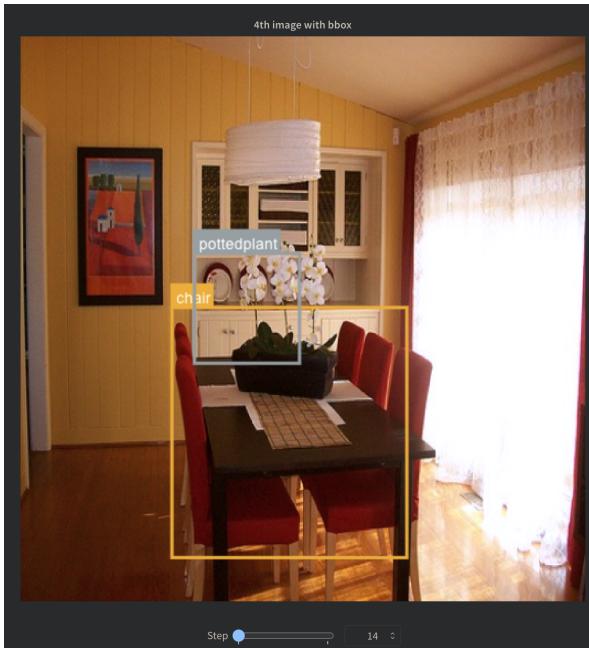
```

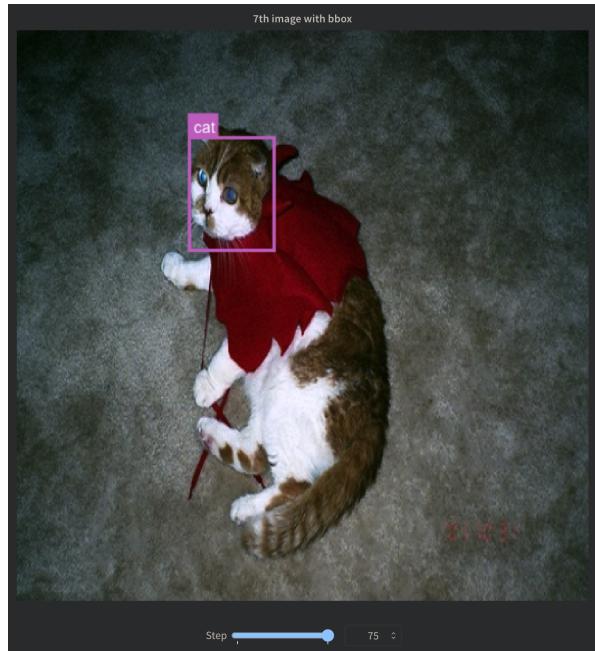
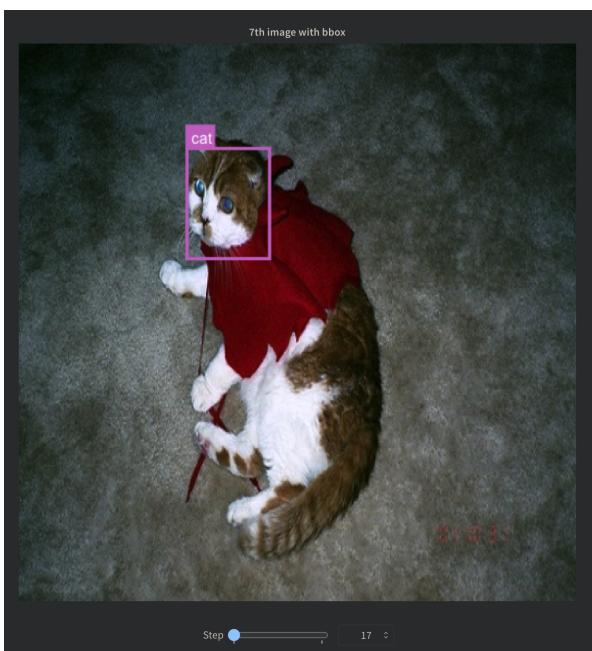
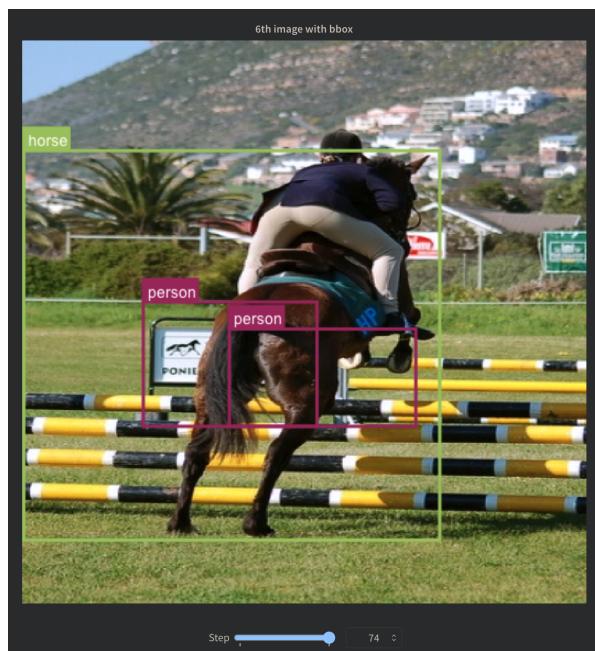
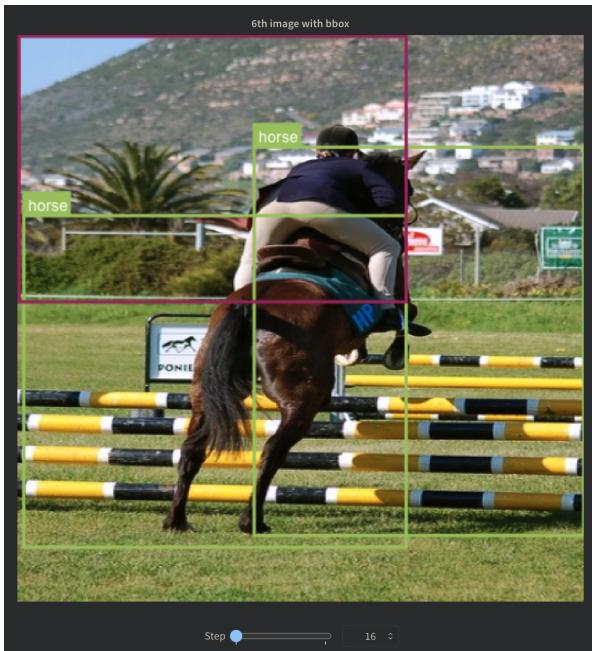
Q2.3

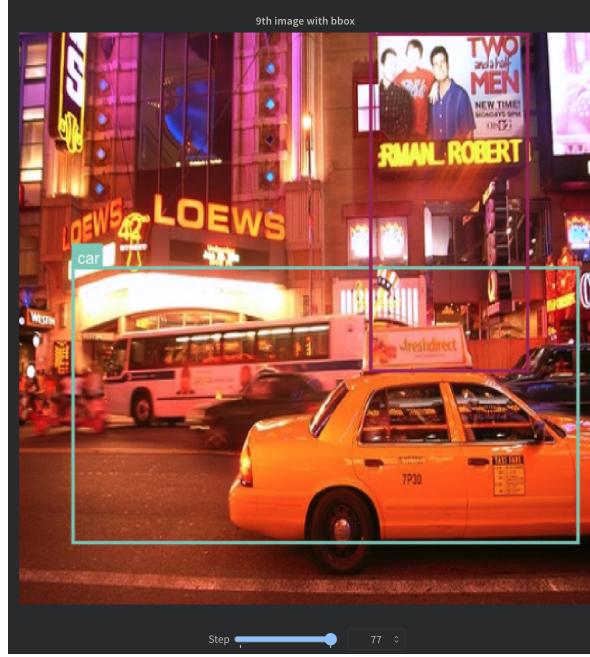
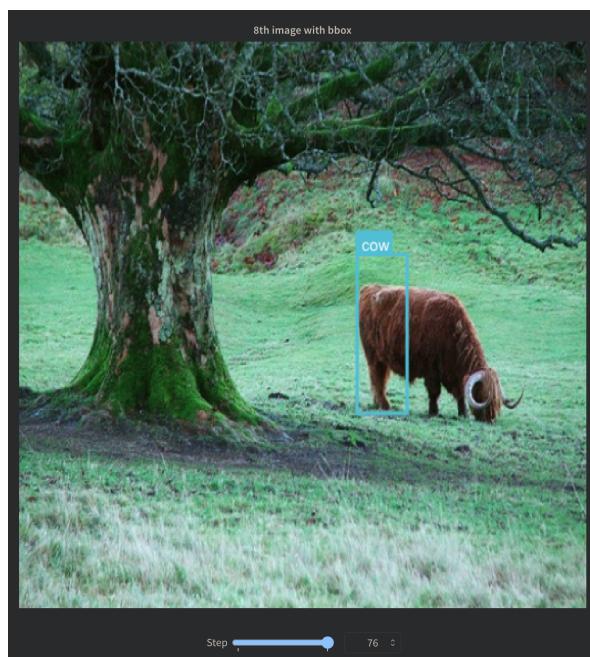
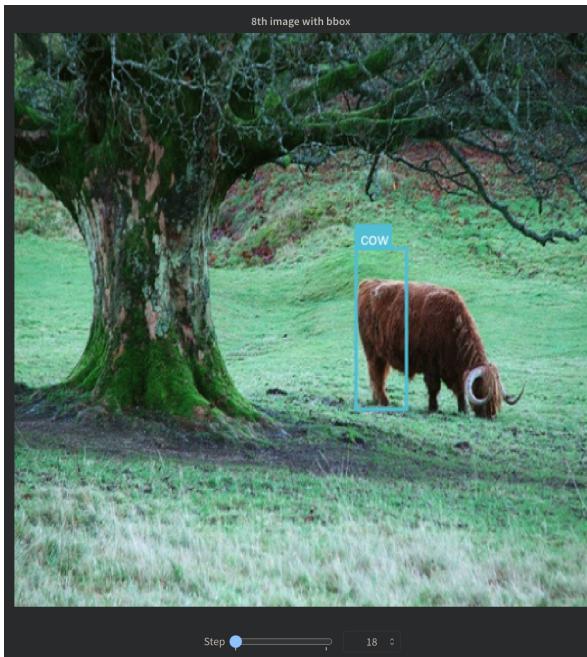












final AP: [0.3114676597298536, 0.26322355466182873, 0.12143539348699098, 0.04246064846296949, 0.0058059822798605994, 0.26113716801143544, 0.21391140685913415, 0.18842382856393336, 0.01045631626372508, 0.08239310206964146, 0.05043815832782442, 0.1562528503884718, 0.28339286269794695, 0.2778879876497678, 0.035609007110676896, 0.04869998075540583, 0.06565027043847324, 0.05325066230271974, 0.3421963485568527, 0.14914270503358934]
final mAP: 0.14816179468255505